

W205, Information Storage and Retrieval

Week #: 9

Lab #: 7

Lab Name: Working with Relational Databases

Name: Gregory Ceccarelli

Date: 10/24/2015

Q1: What is the output of \dt?

Per the terminal stdout, the output is listed below:

List of relations

Schema	Name	Type	Owner
public	actor	table	postgres
public	address	table	postgres
public	category	table	postgres
public	city	table	postgres
public	country	table	postgres
public	customer	table	postgres
public	film	table	postgres
public	film_actor	table	postgres
public	film_category	table	postgres
public	inventory	table	postgres
public	language	table	postgres
public	payment	table	postgres
public	payment_p2007_01	table	postgres
public	payment_p2007_02	table	postgres
public	payment_p2007_03	table	postgres
public	payment_p2007_04	table	postgres
public	payment_p2007_05	table	postgres
public	payment_p2007_06	table	postgres
public	rental	table	postgres
public	staff	table	postgres
public	store	table	postgres

(21 rows)

Q2: What is the schema of the customer table?

Using the \d+ command, the schema and indexes include the below:

Table "public.customer"

Column	Type	Modifiers
customer_id	integer	not null default nextval('customer_customer_id_seq'::regclass)
store_id	smallint	not null
first_name	character varying(45)	not null
last_name	character varying(45)	not null
email	character varying(50)	
address_id	smallint	not null
activebool	boolean	not null default true
create_date	date	not null default ('now'::text)::date
last_update	timestamp without time zone	default now()
active	integer	

Indexes:

- "customer_pkey" PRIMARY KEY, btree (customer_id)
- "idx_fk_address_id" btree (address_id)
- "idx_fk_store_id" btree (store_id)
- "idx_last_name" btree (last_name)

Foreign-key constraints:

- "customer_address_id_fkey" FOREIGN KEY (address_id) REFERENCES address(address_id) ON UPDATE CASCADE ON DELETE RESTRICT
- "customer_store_id_fkey" FOREIGN KEY (store_id) REFERENCES store(store_id) ON UPDATE CASCADE ON DELETE RESTRICT

Q3: What similarities do you see in the explain plans for these 3 queries?

As referenced below, all three query plans suggest the need for a sequential scan.

Query 1 explain plan:

QUERY PLAN

Seq Scan on customer (cost=0.00..14.99 rows=599 width=17)
(1 row)

Query 2 explain plan:

QUERY PLAN

Append (cost=0.00..390.38 rows=4429 width=17)

- > Seq Scan on payment (cost=0.00..0.00 rows=1 width=22)
Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
- > Seq Scan on payment_p2007_01 (cost=0.00..26.36 rows=266 width=16)
Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
- > Seq Scan on payment_p2007_02 (cost=0.00..51.68 rows=531 width=16)
Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
- > Seq Scan on payment_p2007_03 (cost=0.00..126.66 rows=1268 width=16)
Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
- > Seq Scan on payment_p2007_04 (cost=0.00..151.31 rows=1557 width=16)
Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
- > Seq Scan on payment_p2007_05 (cost=0.00..4.73 rows=78 width=15)
Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
- > Seq Scan on payment_p2007_06 (cost=0.00..29.65 rows=728 width=22)
Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))

(15 rows)

Query 3 explain plan:

QUERY PLAN

Append (cost=0.00..390.38 rows=3594 width=12)

- > Seq Scan on payment (cost=0.00..0.00 rows=1 width=18)
Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
- > Seq Scan on payment_p2007_01 (cost=0.00..26.36 rows=242 width=12)
Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
- > Seq Scan on payment_p2007_02 (cost=0.00..51.68 rows=506 width=12)
Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
- > Seq Scan on payment_p2007_03 (cost=0.00..126.66 rows=1290 width=12)
Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
- > Seq Scan on payment_p2007_04 (cost=0.00..151.31 rows=1535 width=12)
Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
- > Seq Scan on payment_p2007_05 (cost=0.00..4.73 rows=13 width=11)
Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
- > Seq Scan on payment_p2007_06 (cost=0.00..29.65 rows=7 width=18)
Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))

(15 rows)

Q4: What is the difference between the plans for the Partitioned table and the union query? Why do you think this difference exists?

As referenced by the query plans below, the difference in the query plans between the partitioned table and the union query is that the union includes two hash aggregates and two group keys. This is because we're combining two physical tables via a union.

Union 2 Tables Query Plan

QUERY PLAN

```
HashAggregate (cost=127.26..129.76 rows=200 width=14)
  Group Key: payment_p2007_01.customer_id
  -> HashAggregate (cost=98.31..109.89 rows=1158 width=26)
    Group Key: payment_p2007_01.payment_id, payment_p2007_01.customer_id,
    payment_p2007_01.staff_id, payment_p2007_01.rental_id, payment_p2007_01.amount,
    payment_p2007_01.payment_date
    -> Append (cost=0.00..80.94 rows=1158 width=26)
      -> Seq Scan on payment_p2007_01 (cost=0.00..23.46 rows=1157 width=26)
        Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
      -> Seq Scan on payment_p2007_02 (cost=0.00..45.90 rows=1 width=26)
        Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
(9 rows)
```

Partitioned Table Query Plan

QUERY PLAN

```
HashAggregate (cost=75.16..77.66 rows=200 width=8)
  Group Key: payment.customer_id
  -> Append (cost=0.00..69.36 rows=1159 width=8)
    -> Seq Scan on payment (cost=0.00..0.00 rows=1 width=14)
      Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
    -> Seq Scan on payment_p2007_01 (cost=0.00..23.46 rows=1157 width=8)
      Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
    -> Seq Scan on payment_p2007_02 (cost=0.00..45.90 rows=1 width=8)
      Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
```

Q5: What join algorithm is used for the inner join?

Per the below referenced query plan, postgres is implementing a hash join:

QUERY PLAN

```
Hash Join (cost=22.48..564.77 rows=17360 width=63)
  Hash Cond: (payment.customer_id = customer.customer_id)
    -> Append (cost=0.00..303.59 rows=17360 width=16)
      -> Seq Scan on payment (cost=0.00..0.00 rows=1 width=22)
      -> Seq Scan on payment_p2007_01 (cost=0.00..20.57 rows=1157 width=16)
      -> Seq Scan on payment_p2007_02 (cost=0.00..40.12 rows=2312 width=16)
      -> Seq Scan on payment_p2007_03 (cost=0.00..98.44 rows=5644 width=16)
      -> Seq Scan on payment_p2007_04 (cost=0.00..117.54 rows=6754 width=16)
      -> Seq Scan on payment_p2007_05 (cost=0.00..3.82 rows=182 width=15)
      -> Seq Scan on payment_p2007_06 (cost=0.00..23.10 rows=1310 width=22)
    -> Hash (cost=14.99..14.99 rows=599 width=49)
      -> Seq Scan on customer (cost=0.00..14.99 rows=599 width=49)
```